

Sardine Refiner

Review & Rubric Quality Guide

Architecture / Interior Design

The Five Rubric Quality Checks

Part 1

1 Coverage

Everything important in the task must be evaluable - including the brief, relevant input information, and supported professional expectations.

2 Overlap

One independent error should not be rewarded or penalized more than once simply because the same issue appears in multiple criteria.

3 Subjectivity

Broad quality language should be translated into observable evidence instead of being ignored or graded only by taste.

The Five Rubric Quality Checks

Part 2

4 **Atomicity**

Each criterion should evaluate one clear requirement or one independent failure mode whenever possible.

5 **Self-Containment**

A reviewer should be able to understand and grade the criterion without guessing what it means or relying on hidden assumptions.

CORE OBJECTIVE

The brief, relevant inputs, outputs, rubric coverage, justifications, PASS/FAIL decisions, pairwise comparison, and final scores should all support the same overall evaluation.

Coverage

Coverage is broader than simply copying the brief into rubrics.

DIRECT-REFERENCE RULE

If the brief says to follow, preserve, reproduce, or apply a provided file directly, the important information contained in that file becomes part of the expected task scope and should be represented in the rubric set.

SPECIFIC-REFERENCE RULE

If the brief uses an input only for a defined purpose - dimensions, layout, openings, material reference, site context, etc. - coverage should focus primarily on those referenced elements rather than every detail in the file.

PROFESSIONAL JUDGMENT

Input files can contain task-critical information that is not repeated word-for-word in the brief. If a detail materially affects correctness, function, safety, coordination, or the ability to deliver the intended work, it may need rubric coverage. Do not promote incidental, legacy, or unrelated file content into mandatory criteria without support.

Coverage - Architecture Examples

How strongly does the brief bind the task to the reference file?

1 **Exact file authority**

Brief: “Follow the provided DWG exactly.” Major footprint, wall layout, openings, circulation cores, levels, and other important DWG-defined geometry may require separate rubric coverage.

2 **Limited file authority**

Brief: “Use the DWG only for room dimensions.” Do not automatically turn unrelated doors, furniture, materials, annotations, or legacy geometry into mandatory requirements.

COVERAGE TEST

Can every important requirement from the brief, the relevant input files, and the professional/domain expectations be evaluated somewhere in the final rubric set?

Coverage - Exact vs Approximate

Do not silently convert flexible language into an exact requirement.

3 **Exact numeric requirement**

Brief: “The ceiling height is 300 cm.” Treat 300 cm as a defined requirement and check the stated value directly, allowing only the measurement or file precision genuinely supported by the source.

4 **Approximate numeric requirement**

Brief: “The ceiling height is approximately 300 cm.” Do not silently convert “approximately” into an exact 300 cm rule. Evaluate reasonable alignment with the intended scale; use a numeric tolerance only when the brief, input, or documented project convention supports it.

KEY DIFFERENCE

Exact language defines a specific target. Approximate language defines an intended range or scale. The rubric should preserve that distinction.

Overlap

Different criteria should evaluate different failure modes - not the same error in different words.

FAST OVERLAP TEST

If this specific error occurs, how many criteria would fail because of exactly the same mistake? If several criteria would fail for the same underlying issue, narrow, merge, or rewrite them.

GOOD SEPARATION

Separate criteria when the conditions can fail independently and each failure has a distinct effect on task quality or compliance.

OVERLAP RISK

A detailed criterion already checks a wall mismatch, while a broader “the model matches the DWG” criterion fails again because of that exact same wall mismatch.

PRACTICAL RULE

Keep one criterion per meaningful failure mode. Broad consistency checks are allowed only when they evaluate a genuinely different level of information.

Borderline Case 1 - Macro-Level Consistency

Useful only when the macro criterion measures a separate proposal-level failure.

WHY THIS IS BORDERLINE

A render-model consistency criterion can be useful because multiple deliverables should represent the same overall proposal. The risk is using that criterion to re-penalize every wall, opening, or material mismatch that already has its own detailed criterion.

- **SAFE:** Use the macro criterion to ask whether the render and model represent the same overall proposal, design version, or architectural concept.
- **DO NOT DOUBLE-PENALIZE:** Do not fail the macro criterion merely because one window, one wall, or one opening already failed a dedicated detailed criterion.
- **FAIL THE MACRO CRITERION:** When the inconsistency is broad enough that the assets no longer represent the same overall proposal - for example, materially different massing, layout logic, facade concept, or project version.
- **WEIGHT CAREFULLY:** A macro criterion should not become a second heavy penalty for a collection of detailed errors unless the combined inconsistency itself creates a separate, meaningful failure.

DECISION RULE

An isolated mismatch belongs to the detailed criterion. A proposal-level contradiction belongs to the macro consistency criterion.

Borderline Case 2 - Reference Detail or Requirement?

Not every visible detail in an input file automatically becomes mandatory.

WHY THIS IS BORDERLINE

Architecture files often contain legacy geometry, alternate options, annotations, furniture, site context, or details that help explain the project but are not necessarily requirements for the current task.

- If the brief says follow or preserve the reference, important reference-defined geometry and organization may become mandatory.
- If the brief says use the file only for a specific purpose, keep the rubric tied to that purpose.
- If the brief explicitly selects certain information from the file, treat those selected elements as direct coverage.
- If the brief is ambiguous, use professional judgment only for details that materially affect the intended result, correctness, coordination, or usability.
- Do not turn incidental annotations, old alternatives, decorative content, or unrelated file information into requirements without support.

DECISION RULE

Ask whether the task would materially change if this reference detail were different. If not, it may be context rather than a rubric requirement.

Subjectivity

Subjective requirements still need coverage - translate them into observable evidence.

DO NOT IGNORE BROAD QUALITY TERMS

Words such as photorealistic, realistic, refined, high-quality, well-proportioned, or visually convincing are still task requirements when the brief uses them.

Example: “Photorealistic rendering” can be broken into:

- Lighting quality and believable light behavior.
- Architectural scale, proportion, and spatial relationships.
- Material realism and surface response.
- Reflections and shadows.
- Realism of secondary elements where relevant.
- Visible geometry, texture, exposure, or rendering defects.

GOAL

Convert broad subjective expectations into observable and defensible evaluation criteria.

Atomicity

Atomicity should prevent both over-penalization and under-penalization.

ONE CLEAR CHECK

Each criterion should evaluate one clear requirement or one independent failure mode whenever possible.

WHY A GENERIC DWG CHECK CAN UNDER-PENALIZE

A criterion such as “The model matches the reference DWG” can hide several independent major errors inside one PASS/FAIL decision. If footprint, wall layout, opening positions, circulation core, and level organization are all wrong, AD1 or AD2 may receive only one penalty even though several major reference failures are present.

BETTER APPROACH

Separate important independent DWG-defined failure modes when they can fail independently: footprint, exterior wall alignment, residential partitions, retail divisions, openings, vertical circulation, levels, or site placement - only where each is relevant.

Atomicity - Avoid Over-Atomization

Atomic does not mean splitting every adjective or every dependent detail into a separate rubric.

WHEN TO COMBINE

If several elements are meaningful only as one requested relationship or combination, evaluate the combination as a whole rather than creating unnecessary micro-criteria.

EXAMPLE

Brief: “The color palette is focused on cool blue, indigo, and dark navy tones.” The requirement is the combined palette. It does not need three separate color-existence criteria unless the task makes those colors independently mandatory.

PRACTICAL RULE

Split independent failure modes. Keep dependent details together when separating them would distort what the brief actually asks for.

Self-Containment

A criterion should be independently understandable and gradable.

SPECIFIC

Avoid vague phrases such as “high resolution,” “sufficient detail,” or “appropriate scale” unless the criterion explains what the reviewer should actually look for.

INDEPENDENT

Do not write criteria that depend on another rubric number or hidden context. Name the element and the expected condition directly.

SUPPORTED

A criterion may rely on the brief, input files, Golden Output, or domain-specific professional expectations - but it should not require undeclared external assumptions.

CONSISTENT VERDICT

The practical target is that different reviewers should reach the same PASS/FAIL decision as often as possible.

Relevance & Weight

Part 1

RELEVANCE TEST

Would the quality, correctness, or compliance of the deliverable meaningfully change depending on whether this criterion passes or fails? If not, reconsider whether the criterion belongs in the set.

8-10 **Most significant criteria**

Critical deliverables, required file types, major required elements, openability, or other failures that strongly affect whether the task is usable or correct.

5-7 **Important criteria**

Meaningful requirements and details such as specific geometry, colors, realism, shapes, or task-specific design conditions.

Relevance & Weight

Part 2

≤5 Minor / primarily aesthetic criteria

Lower-impact details that affect polish but do not strongly determine whether the task is fundamentally correct or usable.

RUBRIC COUNT

Minimum rubric count is a floor, not a target. Complex architecture tasks may legitimately need many more criteria when independent failure modes require separate coverage.

WEIGHT CHECK

The weight should reflect the impact of the criterion itself - not how easy the issue is to notice or how strongly you personally feel about it.

Initial Pairwise Evaluation

Part 1 - Build an independent first impression before the final rubric set.

1 Realism

How believable and professionally produced the deliverable feels.

2 Design and UI/UX Quality

Overall design craft, composition, visual quality, and usability where applicable.

3 Professionalism

Whether the deliverable is polished, free of significant errors or artifacts, and ready for the intended client review.

Initial Pairwise Evaluation

Part 2

4 Coherence and Continuity

Whether the deliverable is internally consistent and logically structured without contradictions.

5 Multi-Asset Consistency

Whether the relevant delivered assets/components share a consistent proposal, style, and quality level.

IMPORTANT

The initial assessment can change after the rubric evidence is complete. The final preference ranking and final rubric scores should tell a consistent story.

Checklist Before Submission

Part 1 - If any answer is NO, revise before submitting.

- 1 Is every criterion atomic?**
- 2 Is every criterion self-contained?**
- 3 Is every criterion relevant?**
- 4 Is everything important covered in the rubrics?**

Checklist Before Submission

Part 2

- 5 Are the rubrics free of duplicates and unnecessary overlap?**
- 6 Is every criterion easy to understand?**
- 7 Do PASS/FAIL verdicts match their justifications?**
- 8 Do the final scores align with the final pairwise ratings and justification?**

Suggested Review Workflow

Part 1 - Validate the task package.

1 READ THE BRIEF FIRST

Check the input file referenced in the brief. Confirm that the actual input matches what the brief describes and decide whether any mismatch prevents reliable completion or evaluation.

2 CHECK THE GOLDEN OUTPUT AGAINST THE BRIEF AND INPUTS

Ask whether the Golden Output could realistically have been produced from this brief and these inputs. Use domain judgment; focus on meaningful reference/alignment problems.

3 CHECK FOR EXTRA OR OUTDATED GOLDEN OUTPUT FILES

Determine whether extra files affect the task. Identify which model/renderers represent the current version and ignore or remove legacy files when appropriate.

4 CHECK THE REQUIRED AD1 / AD2 FILE FORMATS

If the requested format does not work, check whether another usable editable file exists. If no reliable evaluation source remains, the package may be insufficient.

Suggested Review Workflow

Part 2 - Complete the initial comparison and start the rubric build.

5 COMPLETE THE INITIAL PAIRWISE EVALUATIONS

Evaluate RD vs AD1, RD vs AD2, and AD1 vs AD2 before building the final rubric set.

6 CHECK WHETHER RD IS REALLY THE BEST OPTION

If RD does not outperform the alternatives, reconsider the initial judgment. Keep this evidence-based and avoid subjective preference.

7 COPY THE SYNTHETIC RUBRICS INTO A BLANK DOCUMENT

Keep the synthetic set visible as a reference rather than editing it blindly in place.

8 BUILD YOUR OWN SIMPLE RUBRIC OUTLINE FIRST

Use the brief and inputs. Work in category order: Requirements Compliance → Content Correctness → Functionality → Usability & Realism → Presentation & Aesthetics → Editability.

Suggested Review Workflow

Part 3 - Finalize the rubric set and evaluation sources.

9 COMPARE YOUR RUBRICS WITH THE SYNTHETIC SET

Identify missing coverage in both directions. Treat the two sets as references that check each other.

10 CREATE THE FINAL RUBRIC SET

Close important coverage gaps, remove unnecessary overlap, and keep criteria objective, atomic, relevant, and self-contained.

11 ENTER THE FINAL SET INTO THE TASK

Add the criteria in the correct category order.

12 ADD THE CORRECT EVALUATION SOURCE

Specify the appropriate source where needed: 3D model, render, plan, DWG, PDF, material presentation, etc. This helps prevent model/render conflicts.

Suggested Review Workflow

Part 4 - Write justifications and run the first score check.

13 START WRITING THE JUSTIFICATIONS

Keep initial PASS/FAIL observations in a separate document if possible so findings are not lost while editing the rubric set.

14 FLAG SUBJECTIVE JUSTIFICATIONS

If a judgment feels subjective, flag it and return later. Rewrite the criterion or justification around observable evidence where possible.

15 RUN THE FIRST SCORE CHECK

Once the rubric set is complete, click Next and inspect the scores. If RD is unexpectedly low, investigate before continuing.

16 REVIEW RD FAILURES FIRST

Check whether each RD failure is valid, whether the weight is appropriate, and whether the rubric set is representing the task correctly. Do not inflate RD artificially.

Suggested Review Workflow

Part 5 - Validate verdicts and reconcile the final scores.

17 CHECK PASS / FAIL CONSISTENCY

Make sure every justification matches its score. Negative justification + PASS or positive justification + FAIL should trigger a recheck.

18 REVISIT SUBJECTIVE CRITERIA

Make subjective criteria more objective. If a criterion cannot be defended objectively and is not supported by the task, remove it.

19 RUN THE SCORE CHECK AGAIN

Recalculate after all rubric and justification revisions.

20 COMPARE THE FINAL SCORES WITH YOUR INITIAL EVALUATION

Compare the final rubric results with the original RD–AD1, RD–AD2, and AD1–AD2 assessment. They should tell a consistent story. Revise the earlier judgment if the rubric evidence supports a different conclusion.

Suggested Review Workflow

Part 6 - Final uncertainty check and submission.

21 RECHECK EVERYTHING YOU ARE STILL UNSURE ABOUT

Return to any criterion, file, justification, or comparison where doubts remain. Do not leave questionable assumptions unresolved.

22 SUBMIT ONLY WHEN THE TASK FEELS INTERNALLY CONSISTENT

The brief, inputs, outputs, rubric coverage, justifications, PASS/FAIL decisions, pairwise comparison, and final scores should support the same overall evaluation.

WHY THIS WORKS

The same task is checked from several different directions before submission. Consistent use should reduce QC risk, unnecessary regeneration, and review time.